



Chapter 13

Interrupt Architecture, Deferred Procedure Calls, Timers and Thread Scheduling

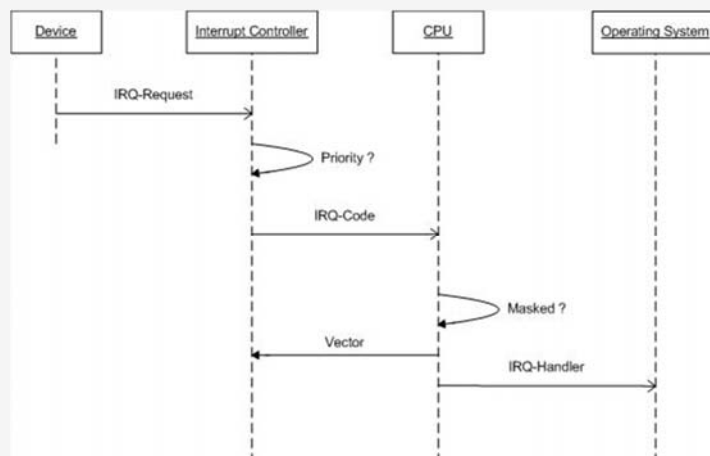
[1]



[2]

Interrupt Processing Sequence

Chapter 13: Interrupts, DPCs, Timers and Thread Scheduling



Interrupt signaling mechanisms

- Line based interrupts
 - Level triggered (e.g. PCI bus)
 - Edge triggered (e.g. legacy ISA bus)
- Message Signaled Interrupt (MSI)
 - Introduced in Vista and above PCI 2.2 spec
 - No sharing / less interrupt latency
 - Enabled by registry subkey of driver's Hardware Key Interrupt Management \ MessageSignaledInterruptProperties
 - 2 phase configuration via `IRP_MJ_PNP`
- MSI-X
 - Defined by PCI 3.0 specification
 - Dynamic configuration by direct call interface to bus driver

Chapter 13: Interrupts, DPCs, Timers and Thread Scheduling

Interrupt priorities and WDM IRQL

- Hardware abstracted WDM IRQLs
- Only three commonly used interrupt levels
 - `PASSIVE_LEVEL`
 - `APC_LEVEL` (rarely used by 3rd party drivers)
 - `DISPATCH_LEVEL`
 - `DIRQL`
- Explicit raising and lowering
 - `KeRaiseIrql(...)` and `KeLowerIrql(...)`
 - Never drop lower than initial IRQL
 - Not usable for synchronization
- DON'T on `IRQL >= DISPATCH_LEVEL`
 - DON'T touch paged memory
 - DON't wait for dispatcher objects

Chapter 13: Interrupts, DPCs, Timers and Thread Scheduling

Interrupt processing

Chapter 13: Interrupts,
DPCs, Timers and Thread
Scheduling

- Interrupt service routine (ISR)
 - Check – is interrupt coming from own device?
 - Communicate with device to stop it from interrupting
- DIRQL interrupt processing
 - Minimize CPU cycles spent on DIRQL
 - Queue DPC (Deferred Procedure Call) for interrupt post processing at DISPATCH_LEVEL
- PASSIVE_LEVEL interrupt processing
 - GPIO secondary interrupts
 - ISR can access devices on SPB at PASSIVE_LEVEL
 - ISR can wait

DPCs and DPC Queues

Chapter 13: Interrupts,
DPCs, Timers and Thread
Scheduling

- DPC Queues (one per processor)
 - DPC object only queued once (per processor)
 - DPCs can run in parallel in multiprocessor environment
- Queue insertion
 - On same processor (by default)
 - Specific processor by:
`KeSetTargetProcessorDpc (...)`
 - End of the queue (by default)
 - Beginning of the queue by:
`KeSetImportanceDpc (...)`



Threaded DPCs

Chapter 13: Interrupts,
DPCs, Timers and Thread
Scheduling

- Supported by Vista and above
- Activation / Deactivation Threaded DPCs
 - Enabled by default
 - HKLM\System\CCS\Control\SessionManager\Kernel\ThreadDpcEnable
 - KeInitializeThreadedDpc(...)
- Programming
 - Executed at IRQL PASSIVE_LEVEL
 - Code must still obey same rules as if running on IRQL DISPATCH_LEVEL



Timers

Chapter 13: Interrupts,
DPCs, Timers and Thread
Scheduling

- Timers DPCs
 - KeInitializeTimer(...)
KeSetTimer(...)
 - 10 milliseconds granularity
- Coalescable timers
 - Available on Windows 7 and above
 - Power saving
 - KeSetCoalescableTimer(...)

Windows Thread Scheduling

Chapter 13: Interrupts,
DPCs, Timers and Thread
Scheduling

Thread execution context

Chapter 13: Interrupts,
DPCs, Timers and Thread
Scheduling

- What makes kernel code run?
- Context of user mode application
 - Request handling only at top level of device stack WDF Driver Object callback
`EvtIoInCallerContext`
 - Driver created system thread in application context
`PsCreateSystemThread(...)`
- System worker threads
 - System thread pool calling e.g. calling
`DriverEntry`
 - Driver created system thread in system context
`PsCreateSystemThread(...)`
- Arbitrary context on elevated IRQL